

Evaluation of Algorithms for Fractal Dimension Estimation in Medical Images: Complexity and Performance Analysis

Gabriel Martins Brum
g.brum@unesp.br

Resumo

Resumo do artigo em português

Palavras-chave: Palavras chave em português

Resumo

Abstract

Resumo do artigo em inglês

Keywords: Palavras chave em inglês

1 Introduction

Fractal geometry, formalised by Benoit B. Mandelbrot in his landmark monograph *The Fractal Geometry of Nature* (MANDELBROT, 1982), provided a mathematical framework for describing the roughness, branching, and statistical self-similarity pervasive in natural phenomena that classical Euclidean geometry fails to capture. The central quantitative descriptor of this framework is the fractal dimension D , estimated in practice as the negative slope of a linear regression over a log-log plot of a covering measure against the measurement scale ε . Shortly after Mandelbrot’s unification, Pentland (PENTLAND, 1984) demonstrated that fractal dimension could be computed directly from digital imagery to discriminate surface textures in natural scenes — a result that bridged fractal geometry and image analysis, and opened the way for its systematic application to medical imaging tasks.

The adoption of fractal analysis in medical imaging occurs since the 1990s, alongside the proliferation of high-resolution digital modalities — multi-detector computed tomography, 3 T magnetic resonance imaging, and digital mammography. Applications spanning pulmonary emphysema (MISHIMA *et al.*, 1999), retinal vascular morphology (JELINEK; FERNANDEZ, 2000), brain cortical folding (FREE *et al.*, 1996), and tumour margin irregularity (RANGAYYAN; NGUYEN, 2007) demonstrated that D encodes structural properties that classical first-order statistical descriptors — variance, entropy, autocorrelation — could not adequately capture. From the 2010s onward, fractal descriptors were incorporated into hybrid feature engineering pipelines alongside learned representations (LOPES; BETROUNI, 2011; KING *et al.*, 2010), and more recently, fractal dimension has been explored as a feature describer to enhance the interpretability of deep learning models (ROBERTO *et al.*, 2021a) and as

a data augmentation strategy with GANs (BORGUE *et al.*, 2025a). Contemporary clinical datasets — commonly comprising images at resolutions exceeding 1024×1024 pixels — elevated the algorithmic efficiency of dimension estimation from a secondary concern to a primary design constraint. The central question shifted from whether fractal dimension is discriminative to at what computational cost it can be reliably extracted, and which algorithm offers the optimal balance of precision, robustness, and efficiency under realistic operating conditions.

Addressing this question requires a rigorous algorithmic engineering perspective. The four methods selected for analysis span the principal design axes of the problem: binary versus grayscale input, and non-overlapping versus dense-overlap covering strategies. The *Classical Box-Counting* (BC) (MANDELBROT, 1982) partitions the binary image into a grid of non-overlapping boxes of decreasing side and counts how many contain foreground pixels, yielding a coarse but efficient estimate of D . The *Gliding Box* (GB) (ALLAIN; CLOITRE, 1991) extends this idea by sliding a window across every pixel position, producing a richer mass-frequency distribution at the cost of processing a far larger number of box placements. The *Differential Box-Counting* (DBC) (SARKAR; CHAUDHURI, 1994) lifts the problem to grayscale images by treating the intensity surface as a three-dimensional relief and counting the column of boxes required to span the local intensity range within each cell, thereby retaining gradient information that binary methods discard. The *Blanket Method* (BM) (PELEG *et al.*, 1984) estimates the fractal dimension of the intensity surface by measuring how its apparent area grows as a morphological cover is progressively inflated and deflated around it at each scale. Contrasting these four methods also surfaces a fundamental trade-off inherent to the pre-processing stage: binarisation lowers the representational

cost of the input, but discards intensity gradients that carry clinically significant discriminatory information — a loss that the grayscale methods avoid at the price of higher computational demand.

This work addresses the following research objectives:

- **Algorithmic analysis:** for each of the four methods, in both naive and structurally optimised formulations, derive formal asymptotic bounds for time and auxiliary space complexity and validate them against empirical measurements on the UCSB medical image dataset, contrasting theoretical scaling predictions with observed wall-clock behaviour.
- **Pre-processing and implementation cost:** quantify the overhead imposed by mandatory pre-processing stages and characterise the data-structure choices that drive the efficiency gap between naive and optimised implementations.
- **Effectiveness and trade-off synthesis:** assess the discriminatory power of each algorithm by measuring the statistical separation it induces between pathological and healthy tissue classes in terms of the estimated fractal dimension, and synthesise all time, memory, and effectiveness evidence into a structured trade-off comparison across the binary-versus-grayscale and naive-versus-optimised design axes.

The remainder of this article is structured as follows: Section 2 establishes the mathematical foundations of fractal analysis and formally defines each algorithm; Section 3 presents pseudocode and derives asymptotic complexity bounds; Section 4 describes the experimental methodology, dataset, and evaluation metrics; Section 5 reports and discusses empirical results; Section 6 concludes with recommendations and directions for future work.

2 Theoretical Foundation

2.1 Fractal Analysis and Dimension Estimation

Fractal geometry provides a mathematical framework for characterising structures whose complexity persists across observation scales, a property termed *statistical self-similarity* (MANDELROT, 1982). Fractal sets exhibit power-law scaling: a characteristic measure $M(\varepsilon)$ obtained at scale ε satisfies

$$M(\varepsilon) \sim C \cdot \varepsilon^{-D}, \quad (1)$$

where C is a proportionality constant and D is the fractal dimension (FALCONER, 2003). Taking logarithms of both sides yields the linear relationship

$$\log M(\varepsilon) = \log C - D \log \varepsilon, \quad (2)$$

from which D is recovered as the negative slope of a least-squares regression over the log-scale pairs $\{(\log \varepsilon_i, \log M(\varepsilon_i))\}_{i=1}^m$, computed across m discrete scales.

Because the Hausdorff dimension (HAUSDORFF, 1919) requires an infimum over all possible coverings, it is not directly computable on discrete pixel grids. The Minkowski–Bouligand (box-counting) dimension D_B provides a tractable approximation (FALCONER, 2003):

$$D_B = \lim_{\varepsilon \rightarrow 0} \frac{\log N(\varepsilon)}{\log(1/\varepsilon)}, \quad (3)$$

where $N(\varepsilon)$ is the minimum number of boxes of side ε needed to cover the set. In practice, Eq. (3) is approximated by the finite regression of Eq. (2) over a scale set $\mathcal{S} = \{\varepsilon_1 > \varepsilon_2 > \dots > \varepsilon_m\}$. The quality of the estimate is assessed through the coefficient of determination R^2 of the regression: values close to unity confirm consistent power-law scaling and are a necessary condition for interpreting D_B as a meaningful fractal descriptor (LOPES; BETROUNI, 2011).

2.2 Classical Box-Counting Method

The Classical Box-Counting (BC) method is the direct discretisation of Eq. (3) for binary images (MANDELROT, 1982; LIEBOVITCH; TOTH, 1989). Let I be a binary image of $N = W \times H$ pixels. For a given box side ε , the image is partitioned into $\lceil W/\varepsilon \rceil \times \lceil H/\varepsilon \rceil$ non-overlapping cells; a cell is *occupied* if it contains at least one foreground pixel:

$$N(\varepsilon) = \sum_{i=0}^{\lceil W/\varepsilon \rceil - 1} \sum_{j=0}^{\lceil H/\varepsilon \rceil - 1} \mathbf{1}[\exists (x, y) \in B_{ij}(\varepsilon) : I(x, y) = 1]. \quad (4)$$

Because BC operates on binary images, a prior binarisation step is required when the input is a grayscale medical image. The standard strategy is Otsu’s global thresholding (OTSU, 1979), which selects the threshold t^* that minimises intra-class intensity variance. This step introduces an inherent trade-off: binarisation lowers the computational cost of the subsequent coverage test, but irreversibly discards intensity gradient information, making the estimated D_B sensitive to the choice of t^* .

To avoid scanning all ε^2 pixels per cell, the optimised formulation pre-computes the *integral image* $P[x][y] = \sum_{i \leq x, j \leq y} I[i][j]$ using the recurrence (CROW, 1984)

$$P[x][y] = I[x][y] + P[x-1][y] + P[x][y-1] - P[x-1][y-1]. \quad (5)$$

Given P , the foreground pixel count within any rectangular cell is retrieved in $\mathcal{O}(1)$ as

$$\begin{aligned} \sigma(x_1, y_1, x_2, y_2) = & P[x_2][y_2] - P[x_1-1][y_2] \\ & - P[x_2][y_1-1] + P[x_1-1][y_1-1], \end{aligned} \quad (6)$$

with occupancy established by $\sigma > 0$. The integral image is built once in $\Theta(N)$ time and space, and shared across all scales in $\mathcal{S}$.

2.3 Gliding Box Method

The Gliding Box (GB) method (ALLAIN; CLOITRE, 1991) replaces the non-overlapping partition of BC with a *dense sliding window*: a box of side L is placed at every valid pixel position (x, y) , and the *box mass*

$$m_{x,y}(L) = \sum_{i=0}^{L-1} \sum_{j=0}^{L-1} I[x+i][y+j] \quad (7)$$

is recorded for each of the $T(L) = (W - L + 1)(H - L + 1)$ valid placements. The resulting *mass-frequency distribution* $n(m, L)$ counts how many placements yield mass m , giving the normalised probability $P(m, L) = n(m, L)/T(L)$.

The fractal dimension is estimated from the mean box mass

$$\mu(L) = \sum_{m=0}^{L^2} m P(m, L), \quad (8)$$

by substituting $\mu(L)$ for $M(\varepsilon)$ in Eq. (2) and fitting the slope over $\mathcal{S}$ (PLOTNICK *et al.*, 1996). A complementary descriptor, *lacunarity*, is obtained from the first two moments of the distribution:

$$\Lambda(L) = \frac{\mu^{(2)}(L) - [\mu(L)]^2}{[\mu(L)]^2}, \quad (9)$$

where $\mu^{(2)}(L) = \sum_m m^2 P(m, L)$. Lacunarity quantifies translational heterogeneity: a homogeneous set yields $\Lambda \approx 0$, while a clustered or gapped structure yields large Λ (ALLAIN; CLOITRE, 1991; TOLLE; MCJUNKIN; GORSICH, 2008). The joint use of D_B and Λ has been shown to improve tissue discriminability in histological image classification (ROBERTO *et al.*, 2021b; BORGUE *et al.*, 2025b), since two textures may share a similar D_B yet differ in their spatial mass distribution.

The same integral-image structure of Eq. (5) applies to GB: pre-computing P reduces each mass evaluation from $\Theta(L^2)$ to $\mathcal{O}(1)$ via Eq. (6), bringing the per-scale cost from $\Theta(NL^2)$ to $\Theta(N)$ (CROW, 1984). This optimisation is more impactful for GB than for BC because GB evaluates $\Theta(N)$ overlapping positions per scale rather than $\Theta(N/L^2)$ disjoint cells, making the naive cost grow as L^2 while the optimised cost remains constant.

2.4 Differential Box-Counting Method

The Differential Box-Counting (DBC) method, proposed by Sarkar and Chaudhuri (SARKAR; CHAUDHURI, 1994), extends BC to grayscale images by treating the intensity surface as a three-dimensional relief. Let I be a grayscale image of $N = W \times H$ pixels with intensities in $\{0, \dots, G-1\}$. For a given scale ε , the image plane is partitioned into $\lceil W/\varepsilon \rceil \times \lceil H/\varepsilon \rceil$ non-overlapping cells as in BC, but each cell is now extended vertically: the intensity axis is divided into bands of height $\delta = \lceil G/\lceil W/\varepsilon \rceil \rceil$, and the number of boxes required to cover the local intensity range within cell (i, j) is

$$n_{ij}(\varepsilon) = \left\lceil \frac{I_{\max}^{ij} - I_{\min}^{ij}}{\delta} \right\rceil + 1, \quad (10)$$

where $I_{\max}^{ij}$ and $I_{\min}^{ij}$ are the maximum and minimum intensities within the cell. The total count $N(\varepsilon) = \sum_{i,j} n_{ij}(\varepsilon)$ is substituted into Eq. (2) to estimate D_B via regression over $\mathcal{S}$.

By retaining gradient information through the intensity range $I_{\max}^{ij} - I_{\min}^{ij}$, DBC avoids the threshold sensitivity inherent to BC, at the cost of processing a richer input representation. The naive formulation scans all ε^2 pixels per cell to compute the local extrema, yielding a per-scale cost of $\Theta(N)$ with a constant proportional to ε^2 . The optimised formulation replaces this scan with two separate *2D range-minimum and range-maximum structures* built in $\Theta(N \log N)$ time and $\Theta(N \log N)$ space via sparse tables, after which each $(I_{\max}^{ij}, I_{\min}^{ij})$ query is answered in $\mathcal{O}(1)$, reducing the per-scale traversal to $\Theta(N/\varepsilon^2)$ with a unit constant — the same asymptotic profile as the optimised BC applied to a grayscale input.

2.5 Blanket Method

The Blanket Method (BM), introduced by Peleg *et al.* (PELEG *et al.*, 1984), estimates the fractal dimension of the intensity surface by measuring how its apparent area grows as a morphological cover is progressively expanded around it at each scale. Let $u_\varepsilon(x, y)$ and $b_\varepsilon(x, y)$ denote the upper and lower *blankets* at dilation level ε , defined by the recurrences

$$u_\varepsilon(x, y) = \max\left(u_{\varepsilon-1}(x, y) + 1, \max_{(s,t) \in \mathcal{N}} u_{\varepsilon-1}(s, t)\right), \quad (11)$$

$$b_\varepsilon(x, y) = \min\left(b_{\varepsilon-1}(x, y) - 1, \min_{(s,t) \in \mathcal{N}} b_{\varepsilon-1}(s, t)\right), \quad (12)$$

with $u_0 = b_0 = I$ and $\mathcal{N}$ denoting the 4-connected neighbourhood of (x, y) . The *blanket volume* at scale ε is

$$V(\varepsilon) = \sum_{x,y} [u_\varepsilon(x, y) - b_\varepsilon(x, y)], \quad (13)$$

and the scale-dependent area measure is obtained as $A(\varepsilon) = [V(\varepsilon) - V(\varepsilon-1)]/2$. Substituting $A(\varepsilon)$ for $M(\varepsilon)$ in Eq. (2) and fitting the slope over $\mathcal{S}$ yields the fractal dimension estimate.

The naive formulation evaluates each recurrence step by scanning the K -neighbourhood of every pixel, incurring $\Theta(NK^2)$ per dilation level and $\Theta(NK^2\varepsilon_{\max})$ in total. The optimised formulation decomposes the 2D neighbourhood maximum and minimum into two successive 1D sliding-window operations — one horizontal, one vertical — each solved in $\Theta(N)$ using a *monotonic deque* (double-ended queue) per row or column (LEMIRE, 2006). This separability reduces the per-level cost to $\Theta(N)$ regardless of neighbourhood size, and the total cost to $\Theta(N\varepsilon_{\max})$. Because BM operates directly on the intensity surface without binarisation, it retains the full gradient information of the image, making it the most information-preserving of the four methods at the price of higher computational demand.

Algorithm 1 Naive Classical Box-Counting

Require: Binary image $I[0 \dots W-1][0 \dots H-1]$, scale set $\mathcal{S}$

Ensure: Pairs $\{(\varepsilon, N(\varepsilon)) : \varepsilon \in \mathcal{S}\}$

```

1: Binarise  $I$  via Otsu's threshold  $t^*$  (OTSU, 1979)
2: for each  $\varepsilon \in \mathcal{S}$  do
3:    $count \leftarrow 0$ 
4:   for  $b_x \leftarrow 0$  to  $\lceil W/\varepsilon \rceil - 1$  do
5:     for  $b_y \leftarrow 0$  to  $\lceil H/\varepsilon \rceil - 1$  do
6:        $occupied \leftarrow \text{FALSE}$ 
7:       for  $dx \leftarrow 0$  to  $\varepsilon - 1$  do
8:         for  $dy \leftarrow 0$  to  $\varepsilon - 1$  do
9:            $px \leftarrow b_x \cdot \varepsilon + dx$ ;  $py \leftarrow b_y \cdot \varepsilon + dy$ 
10:          if  $px < W \wedge py < H \wedge I[px][py] = 1$  then
11:             $occupied \leftarrow \text{TRUE}$ ; break
12:          end if
13:        end for
14:      end for
15:      if  $occupied$  then
16:         $count \leftarrow count + 1$ 
17:      end if
18:    end for
19:  end for
20:  yield  $(\varepsilon, count)$ 
21: end for
22: Estimate  $D_B$  by least-squares regression on  $\{(\log \varepsilon_i, \log N(\varepsilon_i))\}$ 

```

3 Algorithm Project and Analysis

This section presents the algorithmic design and formal asymptotic analysis of the four fractal dimension estimators under study. For each method, a naive brute-force formulation is first described through pseudocode and characterised in terms of its time and auxiliary space complexity; an optimised counterpart is then derived by identifying the computational bottleneck and replacing it with a structurally superior data-structure choice. Throughout this section, $N = W \times H$ denotes the total pixel count of the input image, $m = |\mathcal{S}|$ the number of scales in the evaluation set $\mathcal{S}$, and ε (or L , used interchangeably for the GB) the current box side length. Asymptotic bounds are stated in $\mathcal{O}$, Θ , and Ω notation following the Bachmann–Landau conventions; Θ is used whenever both upper and lower bounds coincide, and $\mathcal{O}$ when only an upper bound is asserted.

3.1 Classical Box-Counting (BC)

3.1.1 Naive Formulation

The naive BC algorithm directly implements Eq. (4) by iterating over every grid cell and scanning all pixels within it to test for occupancy. For each scale $\varepsilon \in \mathcal{S}$, the image domain is partitioned into $\lceil W/\varepsilon \rceil \times \lceil H/\varepsilon \rceil$ non-overlapping cells; within each cell, pixels are visited one by one until a foreground pixel is found or the cell is exhausted (MANDELROT, 1982; LIEBOVITCH; TOTH, 1989). Algorithm 1 details this procedure.

For a fixed ε , the grid contains $\lceil W/\varepsilon \rceil \cdot \lceil H/\varepsilon \rceil =$

Algorithm 2 Optimised Box-Counting (Integral Image)

Require: Binary image $I[0 \dots W-1][0 \dots H-1]$, scale set $\mathcal{S}$

Ensure: Pairs $\{(\varepsilon, N(\varepsilon)) : \varepsilon \in \mathcal{S}\}$

```

1: Binarise  $I$  via Otsu's threshold  $t^*$  (OTSU, 1979)
2: Build integral image  $P$  via Eq. (5)  $\triangleright \Theta(N)$  time and space (CROW, 1984)
3: for each  $\varepsilon \in \mathcal{S}$  do
4:    $count \leftarrow 0$ 
5:   for  $b_x \leftarrow 0$  to  $\lceil W/\varepsilon \rceil - 1$  do
6:     for  $b_y \leftarrow 0$  to  $\lceil H/\varepsilon \rceil - 1$  do
7:        $x_1 \leftarrow b_x \cdot \varepsilon$ ;  $y_1 \leftarrow b_y \cdot \varepsilon$ 
8:        $x_2 \leftarrow \min(x_1 + \varepsilon - 1, W - 1)$ ;  $y_2 \leftarrow \min(y_1 + \varepsilon - 1, H - 1)$ 
9:       if  $\sigma(x_1, y_1, x_2, y_2) > 0$  then  $\triangleright$  Eq. (6),  $\mathcal{O}(1)$ 
10:         $count \leftarrow count + 1$ 
11:      end if
12:    end for
13:  end for
14:  yield  $(\varepsilon, count)$ 
15: end for
16: Estimate  $D_B$  by least-squares regression on  $\{(\log \varepsilon_i, \log N(\varepsilon_i))\}$ 

```

$\Theta(N/\varepsilon^2)$ cells, each requiring at most ε^2 pixel inspections in the worst case (fully foreground cell) and as few as one in the best case (first pixel is foreground). The worst-case per-scale cost is therefore

$$T_{\text{BC,naive}}(\varepsilon) = \Theta\left(\frac{N}{\varepsilon^2}\right) \cdot \mathcal{O}(\varepsilon^2) = \mathcal{O}(N). \quad (14)$$

The lower bound is $\Omega(N/\varepsilon^2)$, since every cell must be visited regardless of its content. Beyond the read-only input buffer of size $\Theta(N)$, the naive formulation requires only a constant number of scalar accumulators, yielding $\Theta(1)$ auxiliary space. Summing over all m scales and including the $\Theta(N)$ binarisation pass:

$$T_{\text{BC,naive}}(N, m) = \Theta(N) + \sum_{\varepsilon \in \mathcal{S}} \mathcal{O}(N) = \mathcal{O}(mN). \quad (15)$$

3.1.2 Optimised Formulation

The inner double loop of Algorithm 1 is the computational bottleneck: it re-inspects up to ε^2 pixels per cell at every scale. This bottleneck is eliminated by pre-computing the two-dimensional prefix sum (integral image) P defined in Eq. (5), which allows any rectangular region sum to be answered in $\mathcal{O}(1)$ using Eq. (6) (CROW, 1984). A cell is occupied if and only if $\sigma > 0$, replacing the entire inner scan with a single arithmetic expression. Algorithm 2 presents the optimised procedure.

The integral image is constructed once in $\Theta(N)$ time and shared across all scales. For each $\varepsilon \in \mathcal{S}$, the outer cell loops visit exactly $\Theta(N/\varepsilon^2)$ cells, each processed in $\mathcal{O}(1)$:

$$T_{\text{BC,opt}}(N, m) = \Theta(N) + \sum_{\varepsilon \in \mathcal{S}} \Theta\left(\frac{N}{\varepsilon^2}\right). \quad (16)$$

Algorithm 3 Naive Gliding Box

Require: Binary image $I[0 \dots W-1][0 \dots H-1]$, scale set $\mathcal{S}$

Ensure: Pairs $\{(L, \mu(L), \Lambda(L)) : L \in \mathcal{S}\}$

- 1: Binarise I via Otsu’s threshold t^* (OTSU, 1979)
- 2: **for** each $L \in \mathcal{S}$ **do**
- 3: $H[0 \dots L^2] \leftarrow 0$ ▷ mass-frequency histogram
- 4: **for** $x \leftarrow 0$ **to** $W - L$ **do**
- 5: **for** $y \leftarrow 0$ **to** $H - L$ **do**
- 6: $mass \leftarrow 0$
- 7: **for** $dx \leftarrow 0$ **to** $L - 1$ **do**
- 8: **for** $dy \leftarrow 0$ **to** $L - 1$ **do**
- 9: $mass \leftarrow mass + I[x+dx][y+dy]$
- 10: **end for**
- 11: **end for**
- 12: $H[mass] \leftarrow H[mass] + 1$
- 13: **end for**
- 14: **end for**
- 15: $T \leftarrow (W-L+1)(H-L+1)$
- 16: Compute $\mu(L)$ and $\mu^{(2)}(L)$ from H and T via Eqs. (8)–(??)
- 17: Compute $\Lambda(L)$ via Eq. (9)
- 18: **yield** $(L, \mu(L), \Lambda(L))$
- 19: **end for**
- 20: Estimate D_B by least-squares regression on $\{(\log L_i, \log \mu(L_i))\}$ (PLOTNICK *et al.*, 1996)

For the typical scale set $\mathcal{S} = \{2^k : k = 1, \dots, \lfloor \log_2 \min(W, H) \rfloor\}$, the sum $\sum_{\varepsilon} N/\varepsilon^2$ is dominated by the smallest scale ($\varepsilon = 2$), giving $\Theta(N/4)$. The total cost is therefore $\Theta(N)$, a strict improvement over the $\mathcal{O}(mN)$ naive bound of Eq. (15). The integral image P requires $\Theta(N)$ auxiliary memory as the sole structure that grows with the input, so the total auxiliary space of the optimised formulation is $\Theta(N)$ — a constant-factor overhead relative to the input. Table 1 consolidates these bounds for both formulations.

Tabela 1: Asymptotic complexity of Classical Box-Counting.

Formulation	Time	Aux. Space
Naive	$\mathcal{O}(mN)$	$\Theta(1)$
Optimised	$\Theta(N)$	$\Theta(N)$

3.2 Gliding Box (GB)

3.2.1 Naive Formulation

The Gliding Box algorithm, introduced by Allain and Cloitre (ALLAIN; CLOITRE, 1991), replaces the disjoint grid of BC with a dense sliding window. For each scale $L \in \mathcal{S}$, the box is positioned at every valid pixel (x, y) , and the box mass $m_{x,y}(L)$ of Eq. (7) is computed by summing all L^2 pixels within the window. The resulting mass-frequency histogram $n(m, L)$ is then normalised to yield $P(m, L)$, from which the mean mass $\mu(L)$ of Eq. (8) and the lacunarity $\Lambda(L)$ of Eq. (9) are computed (PLOTNICK *et al.*, 1996; TOLLE; MCJUNKIN; GORSICH, 2008). Algorithm 3 details the naive procedure.

Algorithm 4 Optimised Gliding Box (Integral Image)

Require: Binary image $I[0 \dots W-1][0 \dots H-1]$, scale set $\mathcal{S}$, integral image P (pre-computed, shared with BC)

Ensure: Pairs $\{(L, \mu(L), \Lambda(L)) : L \in \mathcal{S}\}$

- 1: **for** each $L \in \mathcal{S}$ **do**
- 2: $H[0 \dots L^2] \leftarrow 0$
- 3: **for** $x \leftarrow 0$ **to** $W - L$ **do**
- 4: **for** $y \leftarrow 0$ **to** $H - L$ **do**
- 5: $mass \leftarrow \sigma(x, y, x+L-1, y+L-1)$ ▷ Eq. (6), $\mathcal{O}(1)$ (CROW, 1984)
- 6: $H[mass] \leftarrow H[mass] + 1$
- 7: **end for**
- 8: **end for**
- 9: $T \leftarrow (W-L+1)(H-L+1)$
- 10: Compute $\mu(L)$ and $\mu^{(2)}(L)$ from H and T
- 11: Compute $\Lambda(L)$ via Eq. (9)
- 12: **yield** $(L, \mu(L), \Lambda(L))$
- 13: **end for**
- 14: Estimate D_B by least-squares regression on $\{(\log L_i, \log \mu(L_i))\}$ (PLOTNICK *et al.*, 1996)

For a fixed scale L , there are $T(L) = (W - L + 1)(H - L + 1) = \Theta(N)$ valid window positions, each requiring a full $L \times L$ pixel scan costing $\Theta(L^2)$, so the per-scale cost is

$$T_{\text{GB,naive}}(L) = \Theta(N) \cdot \Theta(L^2) = \Theta(NL^2). \quad (17)$$

The mass-frequency histogram requires $\Theta(L^2)$ entries at scale L , giving $\mathcal{O}(L_{\text{max}}^2)$ auxiliary space — negligible relative to the input when $L_{\text{max}} \ll \sqrt{N}$, but potentially significant for coarse scales. Summing over all m scales and including the binarisation pass:

$$T_{\text{GB,naive}}(N, m) = \Theta(N) + \sum_{L \in \mathcal{S}} \Theta(NL^2) = \Theta\left(N \sum_{L \in \mathcal{S}} L^2\right). \quad (18)$$

For a scale set whose largest element is L_{max} , the sum $\sum L^2$ is $\Theta(mL_{\text{max}}^2)$ in the worst case, making the naive GB substantially more expensive than the naive BC for the same scale set. The dependence on L^2 per position constitutes the algorithmic bottleneck identified for optimisation.

3.2.2 Optimised Formulation

The $\Theta(L^2)$ per-position cost of the naive GB arises entirely from the inner double loop that recomputes the box sum from scratch at each window position. This is identical in structure to the inner scan of naive BC, and is eliminated by the same integral-image mechanism: pre-computing P via Eq. (5) and querying σ via Eq. (6) reduces each mass evaluation to $\mathcal{O}(1)$ (CROW, 1984). Since BC and GB share the same input image, the integral image constructed for BC can be reused directly for GB without additional memory allocation, provided the two methods are executed within the same pipeline. Algorithm 4 presents the optimised procedure.

With $\mathcal{O}(1)$ mass queries, the per-scale cost collapses to the number of window positions:

$$T_{\text{GB,opt}}(L) = \Theta(N) \cdot \mathcal{O}(1) = \Theta(N). \quad (19)$$

Including the single shared $\Theta(N)$ prefix-sum construction and the binarisation pass, the total cost becomes

$$T_{\text{GB,opt}}(N, m) = \Theta(N) + \sum_{L \in S} \Theta(N) = \Theta(mN), \quad (20)$$

a factor of $\Theta(L^2)$ improvement over Eq. (17) at each scale — a gain that grows quadratically with the box side and is therefore most significant at coarse scales. The integral image P is shared with BC and constitutes the sole $\Theta(N)$ auxiliary structure; the per-scale histogram H requires $\Theta(L^2)$ entries, which is typically negligible, keeping the total auxiliary space at $\Theta(N)$. Both optimised formulations therefore share the same $\Theta(N)$ integral image, and their key asymptotic difference lies in the number of positions evaluated per scale: BC visits $\Theta(N/\varepsilon^2)$ disjoint cells, whereas GB visits $\Theta(N)$ overlapping positions regardless of L , incurring a per-scale cost that is a factor of $\Theta(L^2)$ greater than that of optimised BC at the same scale. This overhead is the algorithmic price of the richer mass-frequency distribution that GB produces relative to the binary occupancy count of BC, and which enables the joint estimation of D_B and Λ used in recent histological classification work (ROBERTO *et al.*, 2021b; BORGUE *et al.*, 2025b). Table 2 consolidates the complexity bounds for both GB formulations, and Table 3 contrasts the two methods side by side.

Tabela 2: Asymptotic complexity of the Gliding Box method.

Formulation	Time	Aux. Space
Naive	$\Theta(N \sum_L L^2)$	$\mathcal{O}(L_{\max}^2)$
Optimised	$\Theta(mN)$	$\Theta(N)$

Tabela 3: BC vs. GB optimised formulations (per scale).

Method	Positions/scale	Cost/position	Per-scale cost
BC (opt.)	$\Theta(N/\varepsilon^2)$	$\mathcal{O}(1)$	$\Theta(N/\varepsilon^2)$
GB (opt.)	$\Theta(N)$	$\mathcal{O}(1)$	$\Theta(N)$

4 Methodology

5 Results and Discussion

Resultados da pesquisa

6 Conclusion

Retomada da proposta, resultados e possíveis melhorias/trabalhos futuros.

Tabela 4: Legenda

col1	bol2	bol3	col4
dado1	dado2	dado3	dado4
dado1	dado2	dado3	dado4
dado1	dado2	dado3	dado4
dado1	dado2	dado3	dado4
dado1	dado2	dado3	dado4
resultado1	resultado1	resultado1	resultado1

nota de rodapé

References

- ALLAIN, C.; CLOITRE, M. Characterizing the lacunarity of random and deterministic fractal sets. **Physical Review A**, v. 44, n. 6, p. 3552–3558, 1991.
- BORGUE, V. A. T.; GARCIA, B. L. d. O.; JUNIOR, S. A. P.; ROBERTO, G. F.; ROZENDO, G. B.; NEVES, L. A.; MARTINS, A. S.; TOSTA, T. A. A.; NASCIMENTO, M. Z. d. Association of fractal geometry and data augmentation through gans and xai for classification of histology images. In: **Proceedings**. [S.l.: s.n.], 2025.
- BORGUE, V. A. T.; GARCIA, B. L. d. O.; JÚNIOR, S. A. P.; ROBERTO, G. F.; ROZENDO, G. B.; NEVES, L. A.; MARTINS, A. S.; TOSTA, T. A. A.; NASCIMENTO, M. Z. do. Association of fractal geometry and data augmentation through GANs and XAI for classification of histology images. In: **Proceedings of the 27th International Conference on Enterprise Information Systems (ICEIS 2025)**. [S.l.: s.n.], 2025. v. 1, p. 756–762.
- CROW, F. C. Summed-area tables for texture mapping. In: **Proceedings of the 11th Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH)**. [S.l.: s.n.], 1984. p. 207–212.
- FALCONER, K. **Fractal Geometry: Mathematical Foundations and Applications**. 2. ed. Chichester: John Wiley & Sons, 2003. ISBN 978-0-470-84862-3.
- FREE, S. L.; SISODIYA, S. M.; COOK, M. J.; FISH, D. Z. R.; SHORVON, S. D. Three-dimensional fractal analysis of the white matter surface from magnetic resonance images of the human brain. **Cerebral Cortex**, v. 6, n. 6, p. 830–836, 1996.
- HAUSDORFF, F. Dimension und äußeres maß. **Mathematische Annalen**, v. 79, p. 157–179, 1919. Disponível em: <http://eudml.org/doc/158784>.
- JELINEK, H. F.; FERNANDEZ, E. Neurons and fractals: How reliable and useful are calculations of fractal dimensions? **Journal of Neuroscience Methods**, v. 81, n. 1–2, p. 9–18, 2000.
- KING, R. D.; BROWN, B.; HWANG, M.; JEON, T.; GEORGE, A. T. Fractal dimension analysis of the cortical ribbon in mild Alzheimer’s disease. **NeuroImage**, v. 53, n. 2, p. 471–479, 2010.

- LEMIRE, D. Streaming maximum-minimum filter using no more than three comparisons per element. **Nordic Journal of Computing**, v. 13, n. 4, p. 328–339, 2006. Disponível em: <https://arxiv.org/abs/cs/0610046>.
- LIEBOVITCH, L. S.; TOTH, T. A fast algorithm to determine fractal dimensions by box counting. **Physics Letters A**, v. 141, n. 8–9, p. 386–390, 1989.
- LOPES, R.; BETROUNI, N. Fractal and multifractal analysis: A review. **Medical Image Analysis**, v. 15, n. 4, p. 634–649, 2011.
- MANDELBROT, B. B. **The Fractal Geometry of Nature**. New York: W. H. Freeman, 1982.
- MISHIMA, M.; HIRAI, T.; ITOH, H.; NAKANO, Y.; SAKAI, H.; MURO, S.; NISHIMURA, K.; OKU, Y.; CHIN, K.; OHI, M.; NAKAMURA, T.; BATES, J. H. T.; ALENCAR, A. M.; SUKI, B. Complexity of terminal airspace geometry assessed by fractal analysis. **Journal of Applied Physiology**, v. 87, n. 6, p. 2110–2117, 1999.
- OTSU, N. A threshold selection method from gray-level histograms. **IEEE Transactions on Systems, Man, and Cybernetics**, v. 9, n. 1, p. 62–66, 1979.
- PELEG, S.; NAOR, J.; HARTLEY, R.; AVNIR, D. Multiple resolution texture analysis and classification. **IEEE Transactions on Pattern Analysis and Machine Intelligence**, v. 6, n. 4, p. 518–523, 1984.
- PENTLAND, A. P. Fractal-based description of natural scenes. **IEEE Transactions on Pattern Analysis and Machine Intelligence**, v. 6, n. 6, p. 661–674, 1984.
- PLOTNICK, R. E.; GARDNER, R. H.; HARGROVE, W. W.; PRESTEGAARD, K.; PERLMUTTER, M. Lacunarity analysis: A general technique for the analysis of spatial patterns. **Physical Review E**, v. 53, n. 5, p. 5461–5468, 1996.
- RANGAYYAN, R. M.; NGUYEN, T. M. Fractal analysis of contours of breast masses in mammograms. **Journal of Digital Imaging**, v. 20, n. 3, p. 223–237, 2007.
- ROBERTO, G. F.; LUMINI, A.; NEVES, L. A.; NASCIMENTO, M. Z. do. Fractal neural network: A new ensemble of fractal geometry and convolutional neural networks for the classification of histology images. **Expert Systems with Applications**, Elsevier, v. 166, p. 114103, 2021.
- ROBERTO, G. F.; LUMINI, A.; NEVES, L. A.; NASCIMENTO, M. Z. do. Fractal neural network: A new ensemble of fractal geometry and convolutional neural networks for the classification of histology images. **Expert Systems with Applications**, v. 166, p. 114103, 2021.
- SARKAR, N.; CHAUDHURI, B. B. An efficient differential box-counting approach to compute fractal dimension of image. **IEEE Transactions on Systems, Man, and Cybernetics**, v. 24, n. 1, p. 115–120, 1994.
- TOLLE, C. R.; MCJUNKIN, T. R.; GORSICH, D. J. An efficient implementation of the gliding box lacunarity algorithm. **Physica D: Nonlinear Phenomena**, v. 237, n. 3, p. 306–315, 2008.